


Neural Networks: Backpropagation

+ some SVMs



* Many figures shamelessly taken from Hinton, Karpathy, Efros

Training a neural network




$$O_i = g\left(\sum_j W_{ij} g\left(\sum_k W_{jk} x_k\right)\right)$$



Gradient Descent



Gradient Descent



Gradient Descent

- **Numerical gradient:** easy to write 😊, slow 😞, approximate 😞

- $O(N_w^2)$

$$\frac{df(x)}{dx} = \lim_{h \rightarrow 0} \frac{f(x+h) - f(x)}{h}$$

- **(native) analytic gradient:** exact 😊, complicated 😞, slow 😞

- $O(N_w^2)$

- **back-propagation (cached analytic gradient):** exact 😊, fast 😊, error-prone 😞

- $O(N_w)$, similar to dynamic programming
 - glorified chain rule

In practice: Derive analytic gradient, check your implementation with numerical gradient

The idea behind Backprop

- We don't know what the hidden units ought to do, but we can compute how fast the error changes as we change a hidden activity.
 - Each hidden activity can affect many output units and can therefore have many separate effects on the error. These effects must be combined.
- We can compute error derivatives for all the hidden units efficiently at the same time.
 - Once we have the error derivatives for the hidden activities, it's easy to get the error derivatives for the weights going into a hidden unit.

$$f(x, y) = xy \quad \rightarrow \quad \frac{\partial f}{\partial x} = y \quad \frac{\partial f}{\partial y} = x$$

$$\frac{df(x)}{dx} = \lim_{h \rightarrow 0} \frac{f(x+h) - f(x)}{h}$$

$$f(x+h) = f(x) + h \frac{df(x)}{dx}$$

$$f(x, y) = xy \quad \rightarrow \quad \frac{\partial f}{\partial x} = y \quad \frac{\partial f}{\partial y} = x$$

$$\frac{df(x)}{dx} = \lim_{h \rightarrow 0} \frac{f(x+h) - f(x)}{h}$$

$$f(x+h) = f(x) + h \frac{df(x)}{dx}$$

Example: $x = 4, y = -3. \Rightarrow f(x, y) = -12$

$$\frac{\partial f}{\partial x} = -3 \quad \frac{\partial f}{\partial y} = 4$$

partial derivatives

$$\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$

gradient

Compound expressions: $f(x, y, z) = (x + y)z$

$$q = x + y \quad \frac{\partial q}{\partial x} = 1, \frac{\partial q}{\partial y} = 1$$

$$f = qz \quad \frac{\partial f}{\partial q} = z, \frac{\partial f}{\partial z} = q$$

Compound expressions: $f(x, y, z) = (x + y)z$

$$q = x + y \quad \frac{\partial q}{\partial x} = 1, \frac{\partial q}{\partial y} = 1$$

$$f = qz \quad \frac{\partial f}{\partial q} = z, \frac{\partial f}{\partial z} = q$$

Chain rule:

$$\frac{\partial f}{\partial x} = \frac{\partial f}{\partial q} \frac{\partial q}{\partial x}$$

Compound expressions: $f(x, y, z) = (x + y)z$

$$q = x + y \quad \frac{\partial q}{\partial x} = 1, \frac{\partial q}{\partial y} = 1$$

$$f = qz \quad \frac{\partial f}{\partial q} = z, \frac{\partial f}{\partial z} = q$$

Chain rule:

$$\frac{\partial f}{\partial x} = \frac{\partial f}{\partial q} \frac{\partial q}{\partial x}$$

```
# set some inputs
```

```
x = -2; y = 5; z = -4
```

```
# perform the forward pass
```

```
q = x + y # q becomes 3
```

```
f = q * z # f becomes -12
```

```
# perform the backward pass (backpropagation) in reverse order:
```

```
# first backprop through f = q * z
```

```
dfd_z = q # df/fz = q, so gradient on z becomes 3
```

```
dfd_q = z # df/dq = z, so gradient on q becomes -4
```

```
# now backprop through q = x + y
```

```
dfd_x = 1.0 * dfd_q # dq/dx = 1. And the multiplication here is the chain rule!
```

```
dfd_y = 1.0 * dfd_q # dq/dy = 1
```


Compound expressions: $f(x, y, z) = (x + y)z$

$$q = x + y \quad \frac{\partial q}{\partial x} = 1, \frac{\partial q}{\partial y} = 1$$

$$f = qz \quad \frac{\partial f}{\partial q} = z, \frac{\partial f}{\partial z} = q$$

Chain rule:

$$\frac{\partial f}{\partial x} = \frac{\partial f}{\partial q} \frac{\partial q}{\partial x}$$



```
# set some inputs
```

```
x = -2; y = 5; z = -4
```

```
# perform the forward pass
```

```
q = x + y # q becomes 3
```

```
f = q * z # f becomes -12
```

```
# perform the backward pass (backpropagation) in reverse order:
```

```
# first backprop through f = q * z
```

```
dfd_z = q # df/fz = q, so gradient on z becomes 3
```

```
dfd_q = z # df/dq = z, so gradient on q becomes -4
```

```
# now backprop through q = x + y
```

```
dfd_x = 1.0 * dfdq # dq/dx = 1. And the multiplication here is the chain rule!
```


```
dfd_y = 1.0 * dfdq # dq/dy = 1
```

Another example

$$f(w, x) = \frac{1}{1 + e^{-(w_0 x_0 + w_1 x_1 + w_2)}}$$

Another example:

$$f(w, x) = \frac{1}{1 + e^{-(w_0x_0 + w_1x_1 + w_2)}}$$



$$f(x) = e^x$$

→

$$\frac{df}{dx} = e^x$$

$$f_a(x) = ax$$

→

$$\frac{df}{dx} = a$$

$$f(x) = \frac{1}{x}$$

→


$$\frac{df}{dx} = -1/x^2$$

$$f_c(x) = c + x$$

→

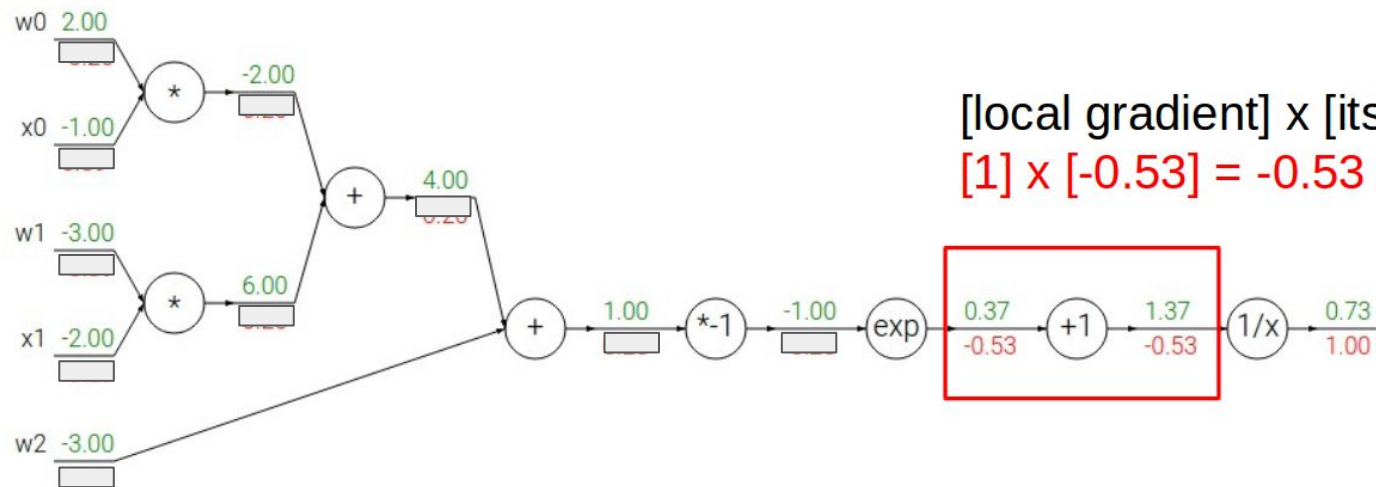
$$\frac{df}{dx} = 1$$

Another example:
$$f(w, x) = \frac{1}{1 + e^{-(w_0x_0 + w_1x_1 + w_2)}}$$




$f(x) = e^x$	\rightarrow	$\frac{df}{dx} = e^x$		$f(x) = \frac{1}{x}$	\rightarrow	$\frac{df}{dx} = -1/x^2$
$f_a(x) = ax$	\rightarrow	$\frac{df}{dx} = a$		$f_c(x) = c + x$	\rightarrow	$\frac{df}{dx} = 1$

Another example: $f(w, x) = \frac{1}{1 + e^{-(w_0x_0 + w_1x_1 + w_2)}}$



$f(x) = e^x$	\rightarrow	$\frac{df}{dx} = e^x$		$f(x) = \frac{1}{x}$	\rightarrow	$\frac{df}{dx} = -1/x^2$
$f_a(x) = ax$	\rightarrow	$\frac{df}{dx} = a$		$f_c(x) = c + x$	\rightarrow	$\frac{df}{dx} = 1$

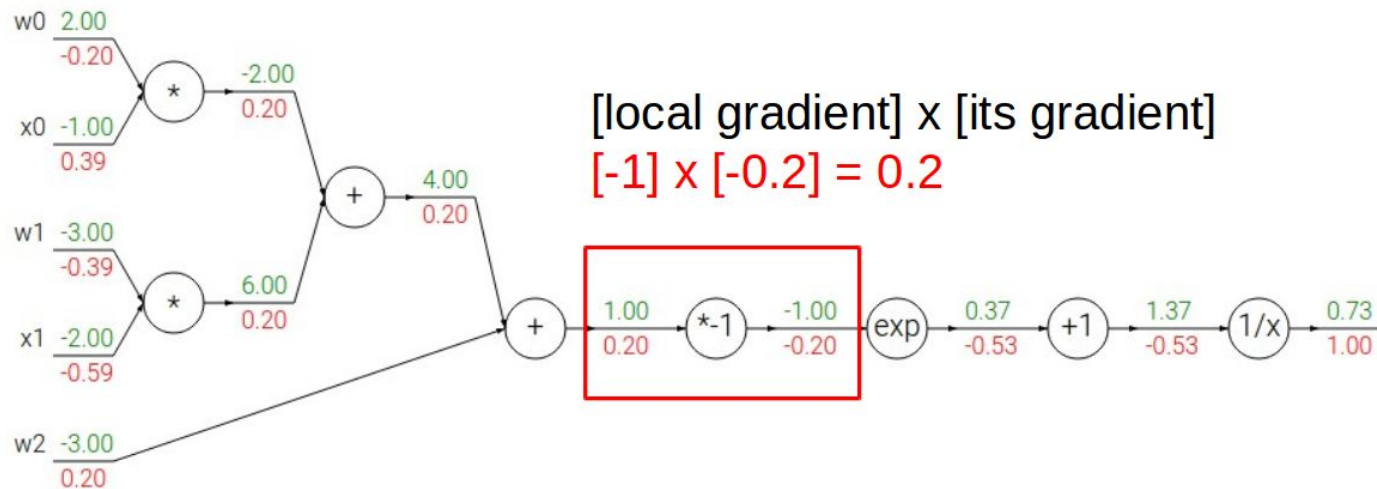
Another example: $f(w, x) = \frac{1}{1 + e^{-(w_0x_0 + w_1x_1 + w_2)}}$



$f(x) = e^x$	\rightarrow	$\frac{df}{dx} = e^x$		$f(x) = \frac{1}{x}$	\rightarrow	$\frac{df}{dx} = -1/x^2$
$f_a(x) = ax$	\rightarrow	$\frac{df}{dx} = a$		$f_c(x) = c + x$	\rightarrow	$\frac{df}{dx} = 1$

Another example:


$$f(w, x) = \frac{1}{1 + e^{-(w_0x_0 + w_1x_1 + w_2)}}$$



$f(x) = e^x$	\rightarrow	$\frac{df}{dx} = e^x$		$f(x) = \frac{1}{x}$	\rightarrow	$\frac{df}{dx} = -1/x^2$
$f_a(x) = ax$	\rightarrow	$\frac{df}{dx} = a$		$f_c(x) = c + x$	\rightarrow	$\frac{df}{dx} = 1$

Another example:

$$f(w, x) = \frac{1}{1 + e^{-(w_0x_0 + w_1x_1 + w_2)}}$$



$$f(x) = e^x$$

→

$$\frac{df}{dx} = e^x$$

$$f_a(x) = ax$$

→

$$\frac{df}{dx} = a$$

$$f(x) = \frac{1}{x}$$

→

$$\frac{df}{dx} = -1/x^2$$


$$f_c(x) = c + x$$

→

$$\frac{df}{dx} = 1$$

Another example:

$$f(w, x) = \frac{1}{1 + e^{-(w_0 x_0 + w_1 x_1 + w_2)}}$$



$f(x) = e^x$	\rightarrow	$\frac{df}{dx} = e^x$		$f(x) = \frac{1}{x}$	\rightarrow	$\frac{df}{dx} = -1/x^2$
$f_a(x) = ax$	\rightarrow	$\frac{df}{dx} = a$		$f_c(x) = c + x$	\rightarrow	$\frac{df}{dx} = 1$

Every gate during backprop computes, for all its inputs:

$$[\text{LOCAL GRADIENT}] \times [\text{GATE GRADIENT}]$$



Can be computed right away,
even during forward pass



The gate receives this during
backpropagation

Backprop: Activation Functions

Logistic (a.k.a
Soft step)



$$f(x) = \frac{1}{1 + e^{-x}}$$

Rectified
Linear Unit
(ReLU)



$$f(x) = \begin{cases} 0 & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$$

Backprop: Activation Functions

Derivatives

Logistic (a.k.a
Soft step)



$$f(x) = \frac{1}{1 + e^{-x}}$$

$$f'(x) = f(x)(1 - f(x))$$

Rectified
Linear Unit
(ReLU)




$$f(x) = \begin{cases} 0 & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$$

$$f'(x) = \begin{cases} 0 & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$$


Before the glory of deep learning: SVMs


- Better theoretical understanding and guarantees
- Works well for many cases
- Simpler model to get working quickly




Support Vector Machines (SVMs)

Find a linear decision surface (**hyperplane**) that can separate patient classes AND has the largest distance (ie largest gap or **margin**) between border-line patients (i.e. **support vectors**)







(a) Training data and an overfitting classifier




(c) Training data and a better classifier




(a) Training data and an overfitting classifier



(b) Applying an overfitting classifier on testing data



(c) Training data and a better classifier



(d) Applying a better classifier on testing data

Lab 4 Tools and Hints

Libraries and Imports:

```
import pandas as pd
from sklearn.neural_network import MLPClassifier
from sklearn.svm import SVC

from sklearn.preprocessing import StandardScaler, MinMaxScaler
from sklearn.preprocessing import LabelEncoder, OneHotEncoder
from sklearn.feature_extraction import DictVectorizer

from sklearn.pipeline import Pipeline
from sklearn.metrics import accuracy_score
from sklearn.model_selection import train_test_split
from sklearn.model_selection import GridSearchCV, ParameterGrid

import numpy as np
```

Lab 4 Tools and Hints:

```
clf = SVC(kernel='linear', random_state=234)

clf.fit(X_train, y_train)
accuracy = clf.score(X_train, y_train)
print('accuracy: {:.3f}%'.format(accuracy*100))
```

```
clf = MLPClassifier(hidden_layer_sizes=(10),
                    activation='logistic',
                    solver='lbfgs',
                    random_state=2563)

clf.fit(X_train, y_train)
accuracy = clf.score(X_train, y_train)
print('accuracy: {:.3f}%'.format(accuracy*100))
```

Lab 4 Tools and Hints:

```
y_pred = clf.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
print('accuracy: {:.3f}%'.format(accuracy*100))
```